

TensorFlow Tutorial

https://docs.google.com/presentation/d/1_H1_XxknAWCwnscVJciN505y7G5hvdSNeF6Ej_3kfYo

git clone

https://github.com/wgchang/POSTECH_CV_Tutorials.git

Contents

1. Tensorflow Basics

- a. Build Computational Graph
- b. Variables
- c. Sharing Variables

2. Image Classification Example

- a. CIFAR-10 classification

Build Computational Graph

Compute '3+4'

```
import tensorflow as tf
```

```
node1 = tf.constant(3.0, tf.float32)
```

```
node2 = tf.constant(4.0) # also tf.float32 implicitly
```

```
print(node1, node2)
```



```
=> Tensor("Const:0", shape=(), dtype=float32) Tensor("Const_1:0", shape=(), dtype=float32)
```

```
sess = tf.Session()
```

```
print(sess.run([node1, node2]))
```

```
=> [3.0, 4.0]
```

sess.run(node) gets the output of the 'node'



Build Computational Graph

Compute '3+4'


```
node3 = tf.add(node1, node2)
```

```
print("node3: ", node3)
```

```
print("sess.run(node3): ", sess.run(node3))
```

```
node3: Tensor("Add_2:0", shape=(), dtype=float32)
```

```
sess.run(node3): 7.0
```



Build Computational Graph

Compute $((3+4)^2)*5$

```
node1 = tf.constant(3.0, tf.float32)
```

```
node2 = tf.constant(4.0, tf.float32)
```

```
node3 = tf.add(node1, node2)
```

```
node4 = tf.constant(2.0, tf.float32)
```


```
node5 = tf.pow(node3, node4)
```

```
node6 = tf.constant(5.0, tf.float32)
```

```
node7 = tf.multiply(node5, node6)
```

```
print(sess.run(node3)) => 7
```

```
print(sess.run(node5)) => 49
```



Build Computational Graph

Compute $((3+4)/(2-5))$

```
node1 = tf.constant(3.0, tf.float32)
```


```
node2 = tf.add(node1, node1)
```

```
node3 = tf.constant(2.0, tf.float32)
```

```
node4 = tf.constant(5.0, tf.float32)
```

```
node5 = tf.subtract(node3, node4)
```

```
node6 = tf.divide(node2, node5)
```



Build Computational Graph

Compute 'x+y'


=> Define placeholder and graph first

=> feed value when sess.run

```
x = tf.placeholder(tf.float32)
```

```
y = tf.placeholder(tf.float32)
```

```
adder = tf.add(x, y) # or just 'x+y' is okay
```



Build Computational Graph

Compute 'x+y'

```
z = sess.run(adder, {x: 3, y: 4.5})
```



```
print(z)
```

```
=> 7.5
```

```
z = sess.run(adder, {x: [1,3], y: [2, 4]})
```

```
print(z)
```

```
=> [ 3.  7.]
```




Build Computational Graph

```
x = tf.placeholder(tf.float32)
y = tf.placeholder(tf.float32)
A = tf.constant([[1., 2.], [3., 4.]])
y_ = tf.matmul(A,x)
diff = y - y_
err = tf.norm(diff)

<=> err = tf.norm(A*x-y)
```

```
print(sess.run(y_, {x: [[1.], [2.]]})) => [[5.], [11.]]
print(sess.run(err, {x: [[1.], [2.]], y:[[8.], [7.]]})) => 5.0
```



Learn Computational Graph

Learn linear model $f(x) \doteq Wx+b$


```
W = tf.Variable([.3], tf.float32) # Variable adds learnable parameters in graph
```

```
b = tf.Variable([-3], tf.float32)
```

```
x = tf.placeholder(tf.float32)
```

```
linear_model = W * x + b
```

```
print(sess.run(linear_model, {x:[1,2,3,4]})) => ERROR!
```



Learn Computational Graph

Learn linear model $f(x) \doteq Wx+b$

```
W = tf.Variable([.3], tf.float32) # Variable adds learnable parameters in graph
```

```
b = tf.Variable([-0.3], tf.float32)
```

```
x = tf.placeholder(tf.float32)
```


```
linear_model = W * x + b
```

```
print(sess.run(linear_model, {x:[1,2,3,4]})) => ERROR!
```

```
init = tf.global_variables_initializer()
```

```
sess.run(init)
```

```
print(sess.run(linear_model, {x:[1,2,3,4]})) => [ 0.      0.30000001 0.60000002 0.90000004]
```



Must initialize variables using initializers!

`global_variables_initializer()` initializes all the variables in a TensorFlow program

Learn Computational Graph

Learn linear model $f(x) \doteq Wx+b$

Loss function : $\sum ||f(x)-y||^2$

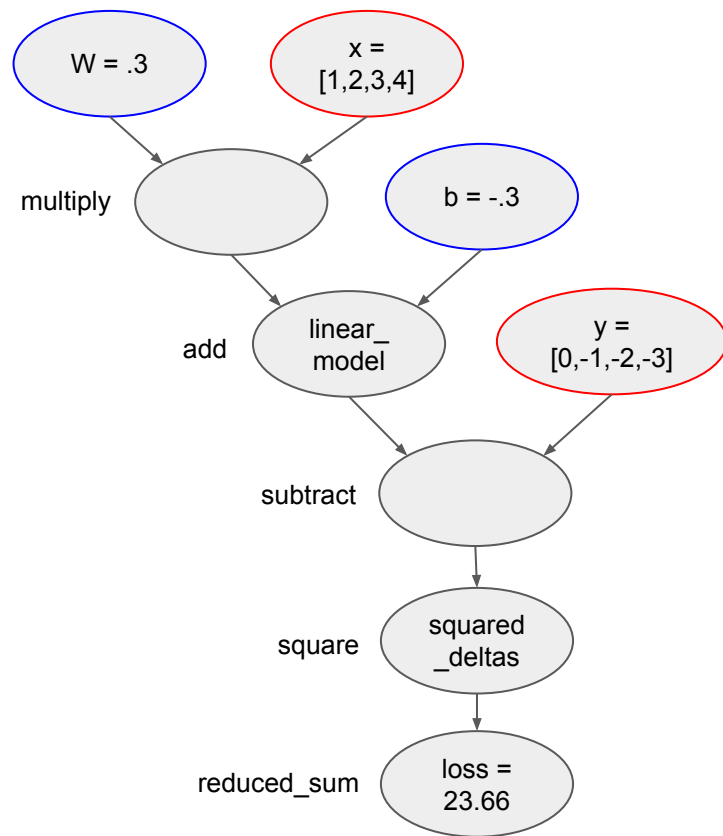
```
y = tf.placeholder(tf.float32)
```

```
squared_deltas = tf.square(linear_model - y)
```

```
loss = tf.reduce_sum(squared_deltas)
```

```
print(sess.run(loss, {x:[1,2,3,4], y:[0,-1,-2,-3]}))
```

23.66



Learn Computational Graph

Learn linear model $f(x) \doteq Wx+b$

Loss function : $\sum ||f(x)-y||^2$

Changing values of Variables

```
fixW = tf.assign(W, [-1.])
```

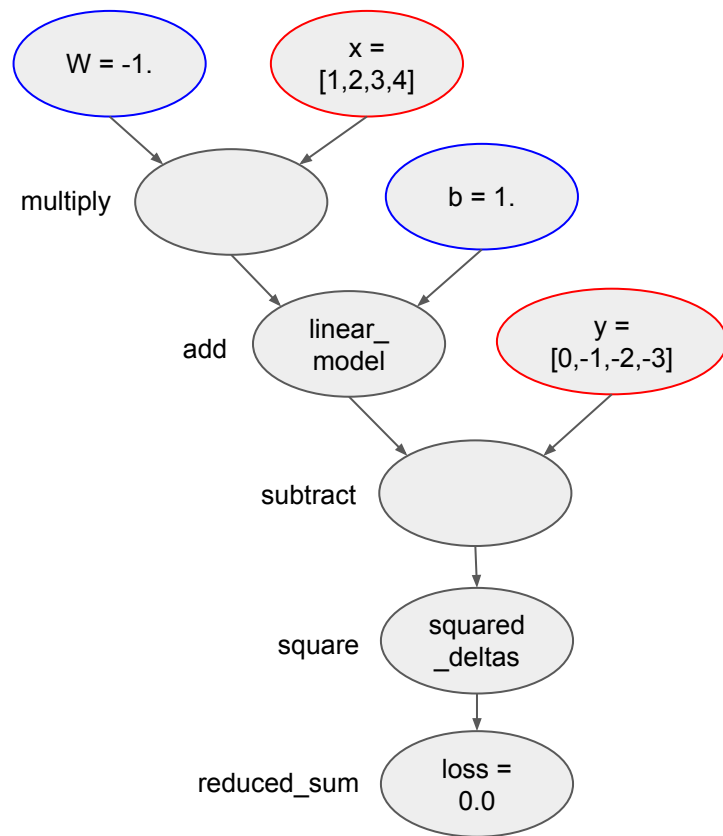
```
fixb = tf.assign(b, [1.])
```

```
sess.run([fixW, fixb])
```

```
print(sess.run(loss, {x:[1,2,3,4], y:[0,-1,-2,-3]}))
```

0.0

Same input & different output!



Learn Computational Graph

Minimize loss function by Optimizer

Simple Example : Stochastic Gradient Descent

```
optimizer = tf.train.GradientDescentOptimizer(0.01)
```

```
train = optimizer.minimize(loss)
```

```
sess.run(init) # reset values to incorrect defaults.
```

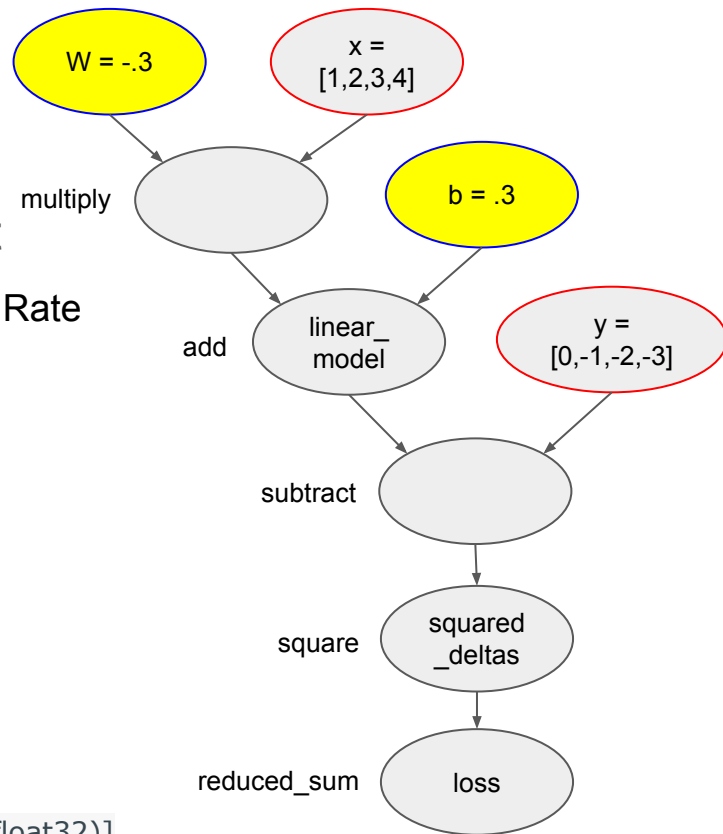
```
for i in range(1000):
```

```
    sess.run(train, {x:[1,2,3,4], y:[0,-1,-2,-3]})
```

```
print(sess.run([W, b]))
```

```
[array([-0.9999969], dtype=float32), array([ 0.99999082], dtype=float32)]
```

Learning Rate



Details about Variables ...

Creation : `tf.Variable(Tensor, name)`

```
weights = tf.Variable(tf.random_normal([784, 200], stddev=0.35), name="weights")
```

size 784 X 200 tensor with random normal values

```
biases = tf.Variable(tf.zeros([200]), name="biases")
```

size 200 tensor with zero values

Set device

Pin a variable to CPU.

```
with tf.device("/cpu:0"):
```

```
    v = tf.Variable(...)
```

Pin a variable to GPU.

```
with tf.device("/gpu:0"):
```

```
    v = tf.Variable(...)
```

Pin a variable to a particular parameter server task.

```
with tf.device("/job:ps/task:7"):
```

```
    v = tf.Variable(...)
```

Operations modifying Variables (`tf.Variable.assign` / `tf.train.Optimizer`) should run on the same device with the mutating Variables

Initializing Variables

`tf.global_variable_initializer()` : An op to initialize ALL the Variables in the model.

```
weights = tf.Variable(tf.random_normal([784, 200], stddev=0.35), name="weights")
```

```
biases = tf.Variable(tf.zeros([200]), name="biases")           # Create two variables.
```

```
...
```

```
init_op = tf.global_variables_initializer()                   # Add an op to initialize the variables.
```

```
# Later, when launching the model
```

```
with tf.Session() as sess:
```

```
    sess.run(init_op)                                         # Run the init operation.
```

```
...
```

```
# Use the model
```

```
...
```

Initializing Variables

If you want to initialize only specific variables :

Use `variables_initialize(var_list)` function, where `var_list` is the set of variables to initialize.

```
W1 = tf.Variable([.2], tf.float32)
```

```
W2 = tf.Variable([.4], tf.float32)
```

```
b = tf.Variable([-0.3], tf.float32)
```

```
init = tf.global_variables_initializer()
```

```
sess.run(init)
```

```
print(sess.run([W1, W2, b]))
```

```
=> [array([ 0.2], dtype=float32), array([ 0.40000001], dtype=float32), array([-0.30000001], dtype=float32)]
```

Initializing Variables

If you want to initialize only specific variables :

Use `variables_initialize(var_list)` function, where `var_list` is the set of variables to initialize.

```
W1 = tf.Variable([.2], tf.float32)      assn_W1 = W1.assign([.1])
W2 = tf.Variable([.4], tf.float32)      assn_W2 = W2.assign([.1])
b = tf.Variable([-.3], tf.float32)      assn_b = b.assign([-.1])

init = tf.global_variables_initializer()  sess.run([assn_W1, assn_W2, assn_b])
sess.run(init)                          print(sess.run([W1, W2, b]))
print(sess.run([W1, W2, b]))
```

```
=> [array([ 0.1], dtype=float32), array([ 0.1], dtype=float32), array([-0.1], dtype=float32)]
```

Initializing Variables

If you want to initialize only specific variables :

Use `variables_initialize(var_list)` function, where `var_list` is the set of variables to initialize.

```
W1 = tf.Variable([.2], tf.float32)
```

```
assn_W1 = W1.assign([.1])
```

```
W2 = tf.Variable([.4], tf.float32)
```

```
assn_W2 = W2.assign([.1])
```

```
b = tf.Variable([-0.3], tf.float32)
```

```
assn_b = b.assign([-0.1])
```

```
init_W = tf.variables_initializer(set([W1, b]))
```

```
sess.run(init_W)
```

```
print(sess.run([W1, W2, b]))
```

```
init = tf.global_variables_initializer()
```

```
sess.run([assn_W1, assn_W2, assn_b])
```

```
sess.run(init)
```

```
print(sess.run([W1, W2, b]))
```

```
print(sess.run([W1, W2, b]))
```

```
=> [array([ 0.2], dtype=float32), array([-0.30000001], dtype=float32)]
```

Initializing Variables

If you want to initialize two Variables A, B to the same values:

Use `initialized_value()` property in the Variable!

```
A = tf.Variable(tf.random_normal([784, 200], stddev=0.35), name="A")
```

Create another variable 'B' with the value of 'A'

```
B = tf.Variable(A.initialized_value(), name="B")
```

Create another variable 'C' with twice the value of 'A'

```
C = tf.Variable(A.initialized_value() * 2.0, name="C")
```

Saving and Loading Variables

Use `tf.train.Saver()`.

You can specify variable name in saving checkpoint file,
or use `tf.Variable.name` as default.

1. Save model

```
# Create some variables.
```

```
v1 = tf.Variable(..., name="v1")
```

```
v2 = tf.Variable(..., name="v2")
```

```
...
```

```
# Add an op to initialize the variables.
```

```
init_op = tf.global_variables_initializer()
```

```
# Add ops to save and restore all the variables.
```

```
saver = tf.train.Saver()
```

```
# Later, launch the model, initialize the variables, do some work
```

```
# , then save the variables to disk.
```

```
with tf.Session() as sess:
```

```
    sess.run(init_op)
```

```
    # Do some work with the model.
```

```
    ..
```

```
    # Save the variables to disk.
```

```
    save_path = saver.save(sess, "model.ckpt")
```

```
    print("Model saved in file: %s" % save_path)
```

Saving and Loading Variables

Use `tf.train.Saver()`.

When you restore variables from a file, you don't have to initialize variables beforehand.

2. Load model

```
# Create some variables.
```

```
v1 = tf.Variable(..., name="v1")
```

```
v2 = tf.Variable(..., name="v2")
```

```
...
```

```
# Add ops to save and restore all the variables.
```

```
saver = tf.train.Saver()
```

```
# Later, launch the model, use the saver to restore variables from disk
```

```
# , and do some work with the model.
```

```
with tf.Session() as sess:
```

```
    # Restore variables from disk.
```

```
    saver.restore(sess, "model.ckpt")
```

```
    print("Model restored.")
```

```
    # Do some work with the model
```

```
    ...
```


Saving and Loading Variables

Save and load only specific variables.

Use python dictionary to pass variables to save/load, key: name, value: variable

```
# Create some variables.
```

```
v1 = tf.Variable(..., name="v1")
```

```
v2 = tf.Variable(..., name="v2")
```

```
...
```

```
# Add ops to save and restore only 'v2' using the name "my_v2"
```

```
saver = tf.train.Saver({"my_v2": v2})
```

```
# Use the saver object normally after that.
```

```
...
```

Saving and Loading Variables

Save and load only specific variables.

```
v1 = tf.Variable(..., name="v1")
```

```
v2 = tf.Variable(..., name="v2")
```

```
...
```

```
# train v1, v2
```

```
saver = tf.train.Saver({"w1": v1, "w2": v2})
```

```
# Use the saver object normally after that.
```

```
save.save(sess, "model.ckpt")
```

```
...
```

```
w1 = tf.Variable(..., name="w1")
```

```
w2 = tf.Variable(..., name="w2")
```

```
w3 = tf.Variable(..., name="w3")
```

```
w4 = tf.Variable(..., name="w3")
```

```
restorer = tf.train.Saver({"w1": w1, "w2": w2})
```

```
# restore w1, w2 then initialize w3, w4
```

```
restorer.restore(sess, "model.ckpt")
```

```
init = tf.variables_initializer(set([w3, w4]))
```

```
sess.run(init)
```

Sharing Variables

Imagine you created image filters with 2 convolutional layers this way:


```
def my_image_filter(input_images):  
    conv1_weights = tf.Variable(tf.random_normal([5, 5, 32, 32]), name="conv1_weights")  
    conv1_biases = tf.Variable(tf.zeros([32]), name="conv1_biases")  
    conv1 = tf.nn.conv2d(input_images, conv1_weights, strides=[1, 1, 1, 1], padding='SAME')  
    relu1 = tf.nn.relu(conv1 + conv1_biases)  
  
    conv2_weights = tf.Variable(tf.random_normal([5, 5, 32, 32]), name="conv2_weights")  
    conv2_biases = tf.Variable(tf.zeros([32]), name="conv2_biases")  
    conv2 = tf.nn.conv2d(relu1, conv2_weights, strides=[1, 1, 1, 1], padding='SAME')  
    return tf.nn.relu(conv2 + conv2_biases)
```

Sharing Variables

Problem: You want to process two images (image1, image2) to same filters but whenever you call 'my_image_filter', new variables are created

```
# First call creates one set of 4 variables.
```

```
result1 = my_image_filter(image1)
```



Sharing Variables

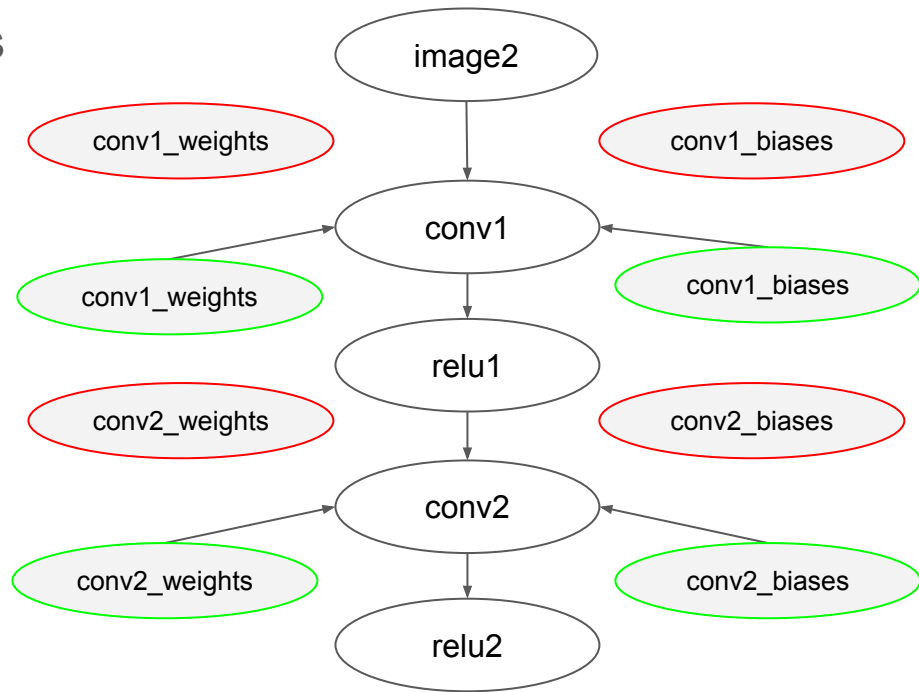
Problem: You want to process two images (image1, image2) to same filters but whenever you call 'my_image_filter', new variables are created

```
# First call creates one set of 4 variables.
```

```
result1 = my_image_filter(image1)
```

```
# Another set of 4 variables is created in the second call.
```

```
result2 = my_image_filter(image2)
```



Sharing Variables

One way to solve the problem... : Use separate code for filters (ex. dictionary)

```
variables_dict = { "conv1_weights": tf.Variable(tf.random_normal([5, 5, 32, 32]), name="conv1_weights")
                  "conv1_biases": tf.Variable(tf.zeros([32]), name="conv1_biases")
                  ... etc. ...}
```

```
def my_image_filter(input_images, variables_dict):
    conv1 = tf.nn.conv2d(input_images, variables_dict["conv1_weights"], strides=[1, 1, 1, 1], padding='SAME')
    relu1 = tf.nn.relu(conv1 + variables_dict["conv1_biases"])
    conv2 = tf.nn.conv2d(relu1, variables_dict["conv2_weights"], strides=[1, 1, 1, 1], padding='SAME')
    return tf.nn.relu(conv2 + variables_dict["conv2_biases"])
```

```
result1 = my_image_filter(image1, variables_dict)
result2 = my_image_filter(image2, variables_dict) # Both calls to my_image_filter() now use the same variables
```

Sharing Variables

This is BAD solution... since,

Breaks Encapsulation!!

- The code that builds the graph must document the names, types, and shapes of variables to create.
- When the code changes, the callers may have to create more, or less, or different variables.

Solution : Use **tf.get_variable()** and **tf.variable_scope()**

Sharing Variables

`tf.get_variable(<name>, <shape>, <initializer>)`

- Get or create a variable instead of `tf.Variable()`
- Case I : Create a variable if a variable named with `<name>` doesn't exist, and raise error if does exist.
- Case II : Raise error if a variable named with `<name>` doesn't exist, but if does exists, then return that variable.

`tf.variable_scope(<scope_name>)`

- adds `<scope_name>` as the prefix of the `<name>` passed to `get_variable`
- Decide whether to reuse variables inside the scope by 'reuse' flag.
 - `reuse == False => Case I.` `reuse == True => Case II.`

Sharing Variables

Some initializers in `get_variable...`

- `tf.constant_initializer(value)` : initializes everything to the provided value
- `tf.random_uniform_initializer(a, b)` : initializes uniformly from `[a, b]`
- `tf.random_normal_initializer(mean, stddev)` : initializes from the normal distribution with the given mean and standard deviation.

Sharing Variables

Adds <scope_name> in front of the variable name <name>

```
with tf.variable_scope("foo"):
```

```
    v = tf.get_variable("v", [1])
```

```
assert v.name == "foo/v:0"
```

```
with tf.variable_scope("foo"):
```

```
    with tf.variable_scope("bar"):
```

```
        v = tf.get_variable("v", [1])
```

```
assert v.name == "foo/bar/v:0"
```

By setting 'reuse' flag of variable_scope true, we can reuse variables by get_variable. The default flag is False.

```
with tf.variable_scope("foo"):
```

```
    v = tf.get_variable("v", [1])
```

```
with tf.variable_scope("foo", reuse=True):
```

```
    v1 = tf.get_variable("v", [1])
```

```
assert v1 is v      # v1 = v.
```

Sharing Variables

`tf.get_variable_scope()` retrieves the current variable scope and by using `reuse_variables()`, you can change the reuse flag to True to reuse variables inside the scope.

```
with tf.variable_scope("foo"):
    v = tf.get_variable("v", [1])
    tf.get_variable_scope().reuse_variables()
    v1 = tf.get_variable("v", [1])
assert v1 is v
```

tf.get_variable_scope().reuse == False

tf.get_variable_scope().reuse == True

Setting 'reuse' back to False is impossible! => Breaking contract inside the scope.

Sharing Variables

Enter inside the reusing
variable scope,

then exit back to non-reusing
one.

```
with tf.variable_scope("root"):
    # At start, the scope is not reusing.
    assert tf.get_variable_scope().reuse == False
    with tf.variable_scope("foo"):
        # Opened a sub-scope, still not reusing.
        assert tf.get_variable_scope().reuse == False
    with tf.variable_scope("foo", reuse=True):
        # Explicitly opened a reusing scope.
        assert tf.get_variable_scope().reuse == True
        with tf.variable_scope("bar"):
            # Now sub-scope inherits the reuse flag.
            assert tf.get_variable_scope().reuse == True
        # Exited the reusing scope, back to a non-reusing one.
    assert tf.get_variable_scope().reuse == False
```

Sharing Variables

Alternative way... : Capturing Variable Scope

Can reuse variables without changing scope name

```
with tf.variable_scope("foo") as foo_scope:
```

```
    v = tf.get_variable("v", [1])
```

```
with tf.variable_scope(foo_scope):
```

```
    w = tf.get_variable("w", [1])
```

```
with tf.variable_scope(foo_scope, reuse=True):
```

```
    v1 = tf.get_variable("v", [1])
```

```
    w1 = tf.get_variable("w", [1])
```

```
assert v1 is v
```

```
assert w1 is w
```

```
# Doesn't change scope name even though it is nested inside.
```

```
with tf.variable_scope("foo") as foo_scope:
```

```
    assert foo_scope.name == "foo"
```

```
with tf.variable_scope("bar"):
```

```
    with tf.variable_scope("baz") as other_scope:
```

```
        assert other_scope.name == "bar/baz"
```

```
    with tf.variable_scope(foo_scope) as foo_scope2:
```

```
        assert foo_scope2.name == "foo" # Not changed.
```

Sharing Variables

How to solve the previous problem?

```
def conv_relu(input, kernel_shape, bias_shape):  
    # Create variable named "weights", "biases".  
    weights = tf.get_variable("weights", kernel_shape, initializer=tf.random_normal_initializer())  
    biases = tf.get_variable("biases", bias_shape, initializer=tf.constant_initializer(0.0))  
    conv = tf.nn.conv2d(input, weights, strides=[1, 1, 1, 1], padding='SAME')  
    return tf.nn.relu(conv + biases)  
  
def my_image_filter(input_images):  
    with tf.variable_scope("conv1"):  
        # Variables created here will be named "conv1/weights", "conv1/biases".  
        relu1 = conv_relu(input_images, [5, 5, 32, 32], [32])  
    with tf.variable_scope("conv2"):  
        # Variables created here will be named "conv2/weights", "conv2/biases".  
        return conv_relu(relu1, [5, 5, 32, 32], [32])
```

Sharing Variables


```
result1 = my_image_filter(image1)
```

```
result2 = my_image_filter(image2)
```

```
# Raises ValueError(... conv1/weights already exists ...)
```

```
with tf.variable_scope("image_filters") as scope:
```

```
    result1 = my_image_filter(image1)
```



Sharing Variables

```
result1 = my_image_filter(image1)
```

```
result2 = my_image_filter(image2)
```


```
# Raises ValueError(... conv1/weights already exists ...)
```

```
with tf.variable_scope("image_filters") as scope:
```

```
    result1 = my_image_filter(image1)
```

```
    scope.reuse_variables()
```

```
    result2 = my_image_filter(image2)
```



More about variable scope...

`tf.variable_scope()` can carry the **default initializer** for `get_variable`!

```
with tf.variable_scope("foo", initializer=tf.constant_initializer(0.4)):
    v = tf.get_variable("v", [1])
    assert v.eval() == 0.4 # Default initializer as set above.
    w = tf.get_variable("w", [1], initializer=tf.constant_initializer(0.3)):
    assert w.eval() == 0.3 # Specific initializer overrides the default.
    with tf.variable_scope("bar"):
        v = tf.get_variable("v", [1])
        assert v.eval() == 0.4 # Inherited default initializer.
    with tf.variable_scope("baz", initializer=tf.random_normal_initializer()):
        v = tf.get_variable("v", [1])
        assert v.eval() == 0.176 # Changed default initializer.
```

Name scope

`tf.name_scope(<scope_name>)` governs operation name only,

while `variable_scope` governs both operation and `tf.Variable` name, but not `tf.get_variable`

```
with tf.variable_scope("foo"):
    x = 1.0 + tf.get_variable("v", [1])
assert x.op.name == "foo/add"
```

```
with tf.variable_scope("foo"):
    with tf.name_scope("bar"):
        v = tf.get_variable("v", [1])
        w = tf.Variable([.1], tf.float32)
        x = 1.0 + v
    assert v.name == "foo/v:0"
    assert w.name == "foo/bar/w:0"
    assert x.op.name == "foo/bar/add"
```

Reference

Build Computational Graph

https://www.tensorflow.org/get_started/get_started

Variables

https://www.tensorflow.org/programmers_guide/variables

Sharing Variables

https://www.tensorflow.org/programmers_guide/variable_scope