

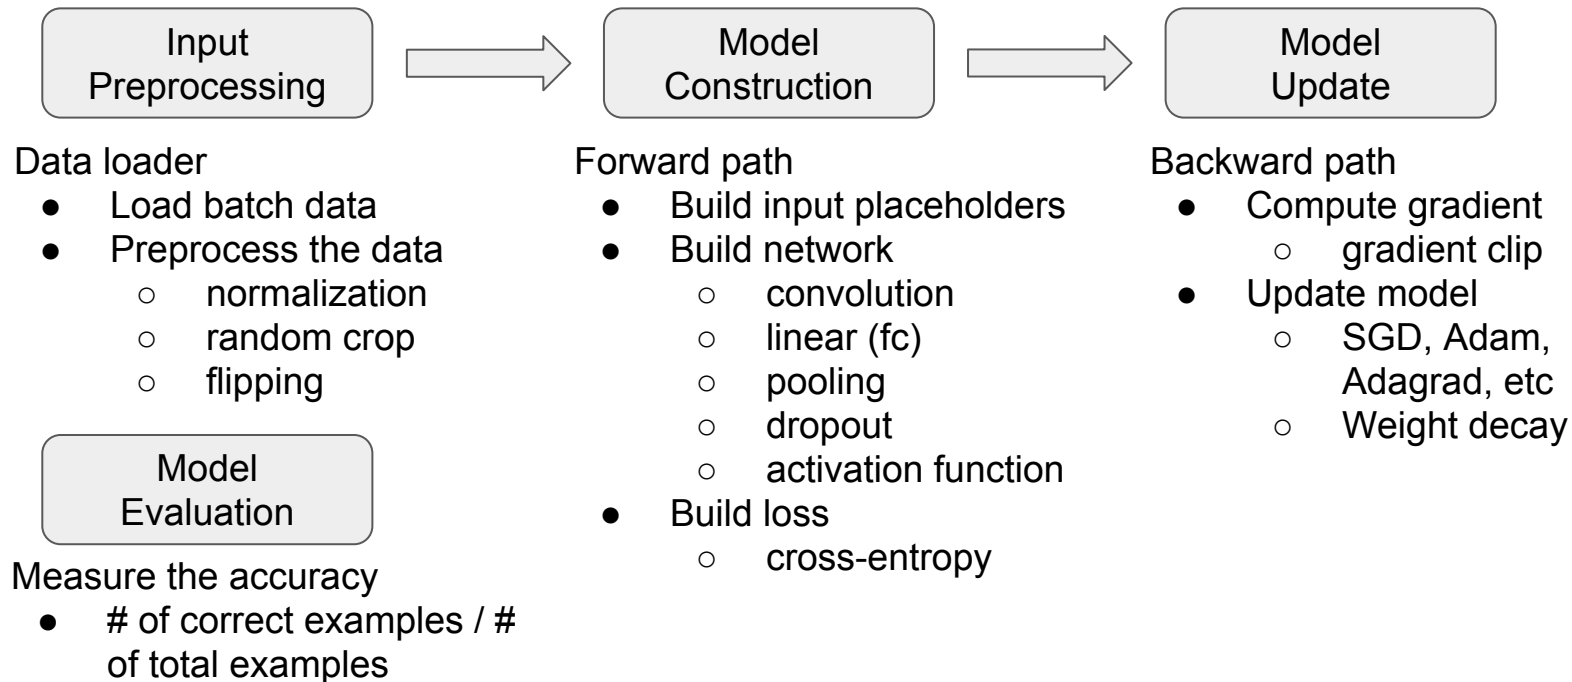
Image Classification

Image Classification

Given an image, identifying which object is in the image

			
mite	container ship	motor scooter	leopard
<div> <div></div> <div>mite</div> <div>black widow</div> <div>cockroach</div> <div>tick</div> <div>starfish</div> </div>	<div> <div></div> <div>container ship</div> <div>lifeboat</div> <div>amphibian</div> <div>fireboat</div> <div>drilling platform</div> </div>	<div> <div></div> <div>motor scooter</div> <div>go-kart</div> <div>moped</div> <div>bumper car</div> <div>golfcart</div> </div>	<div> <div></div> <div>leopard</div> <div>jaguar</div> <div>cheetah</div> <div>snow leopard</div> <div>Egyptian cat</div> </div>
			
grille	mushroom	cherry	Madagascar cat
<div> <div></div> <div>convertible</div> <div>grille</div> <div>pickup</div> <div>beach wagon</div> <div>fire engine</div> </div>	<div> <div></div> <div>agaric</div> <div>mushroom</div> <div>jelly fungus</div> <div>gill fungus</div> <div>dead-man's-fingers</div> </div>	<div> <div></div> <div>dalmatian</div> <div>grape</div> <div>elderberry</div> <div>ffordshire bullterrier</div> <div>currant</div> </div>	<div> <div></div> <div>squirrel monkey</div> <div>spider monkey</div> <div>titi</div> <div>indri</div> <div>howler monkey</div> </div>

Image Classification



Input Preprocessing

CIFAR-10 Dataset

- **10 classes** - airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck
- **32x32 color images** with 6,000 images per class (50,000 training / 10,000 test)

airplane



automobile



bird



cat



deer



dog




frog



horse



ship



truck



Input Preprocessing


Data loader (cifar10_loader.py) does

- Reshape a given vector of 3,072 size into a tensor of 3x32x32 size
 - (channel, width, height)
- Transpose (channel, width, height) into (width, height, channel)
- Load data of batch size using `get_batch()`

Download and check the dataset in the ipython notebook
(002_image_classification/image_classification.ipynb)

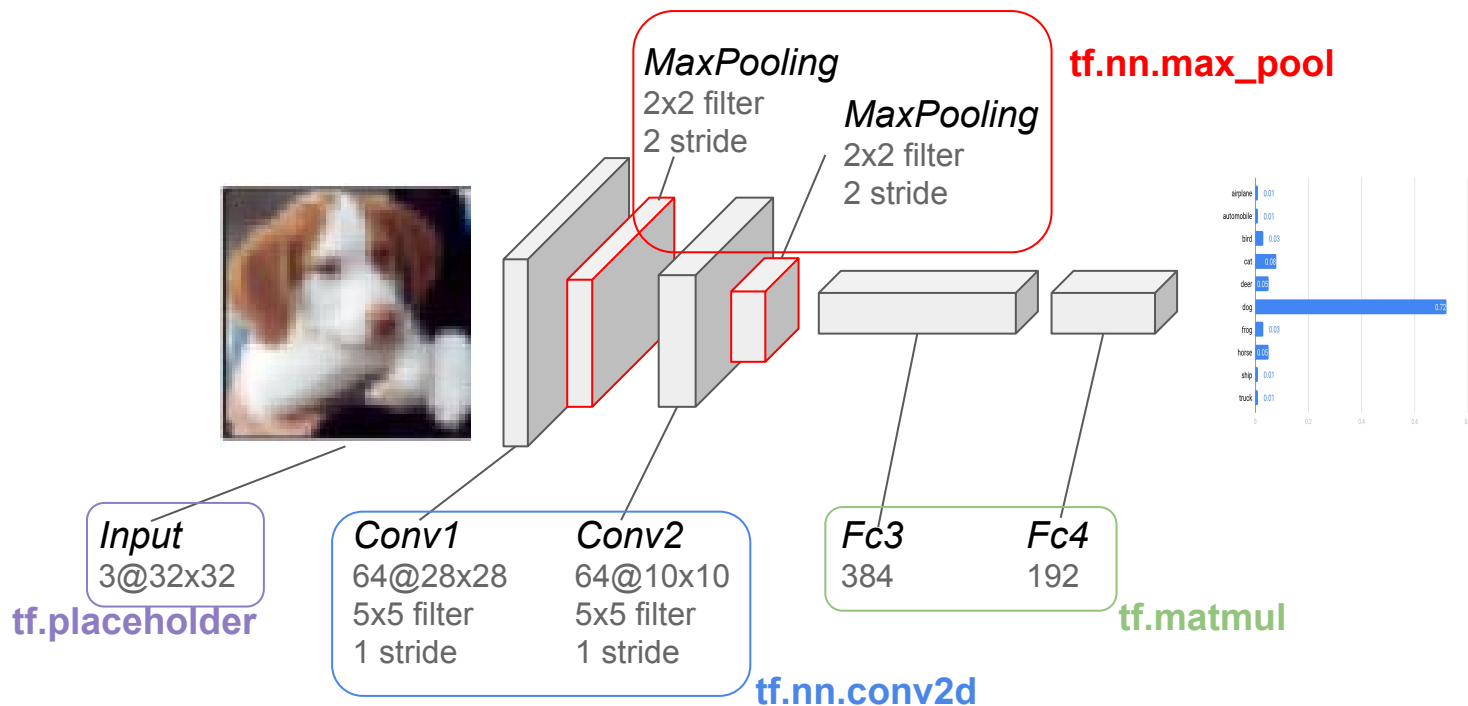
Model Construction

2 convolutional layers + 2 fully connected layers + classification layer



Model Construction

How to implement the layers in TensorFlow?



Model Construction

How to implement the layers?

Convolutional Layer - `tf.nn.conv2d()`

```
# Create variables
filter = tf.get_variable("weights", [filter_size, filter_size,
inp_dim, out_dim])
bias = tf.get_variable("biases", [out_dim])

# Compute convolution
conv = tf.nn.conv2d(input, filter)
conv = tf.nn.bias_add(conv, bias)
```

Fully Connected Layer - `tf.matmul()`

```
# Create variables
weight = tf.get_variable("weights", [inp_dim, out_dim])
bias = tf.get_variable("biases", [out_dim])

# Compute convolution
fc = tf.matmul(input, weight)
fc = tf.nn.bias_add(fc, bias)
```

If you want to apply activation function (relu, tanh)?

```
conv = tf.nn.relu(conv)
conv = tf.tanh(conv)
```

```
fc = tf.nn.relu(fc)
fc = tf.tanh(fc)
```

Max pooling?

```
# kernel size 2 and stride 2
pool = tf.nn.max_pool(conv, ksize=[1,2,2,1],
strides=[1,2,2,1])
```

Dropout?

```
drop_fc = tf.nn.dropout(fc, keep_prob)
```


Model Construction

We should always define variables manually? **No**

There are three libraries based on TensorFlow which are **simple and concise!**

```
# Create variables
filter = tf.get_variable("weights", [filter_size, filter_size,
inp_dim, out_dim])
bias = tf.get_variable("biases", [out_dim])

# Compute convolution
conv = tf.nn.conv2d(input, filter)
conv = tf.nn.bias_add(conv, bias)

# Apply activation function
conv = tf.nn.relu(conv)
conv = tf.tanh(conv)

# Apply dropout
drop_conv = tf.nn.dropout(conv, keep_prob)
```

→
wrapper
function


1. TFLearn
 - `tflearn.layers.conv.conv_2d()`
2. Keras
 - `tf.contrib.keras.layers.Conv2D()`
3. TF-Slim
 - `tf.contrib.slim.conv2d()`

Note that we do not cover above libraries in this course

Model Update

Define loss function (cross-entropy loss)

Define optimizer (Adam)



Model Update

How to define loss function and optimizer in TensorFlow?

Loss Function

Compute cross-entropy loss

```
cross_entropy =  
tf.nn.sparse_softmax_cross_entropy_with_logits(  
labels, logits)  
loss = tf.reduce_mean(cross_entropy)
```

* cross-entropy loss?

- $\text{LogSoftMax} + \text{ClassNLLCriterion}$
- $\text{loss}(x, \text{class})$
 $= -\log(\exp(x[\text{class}]) / \sum_j \exp(x[j]))$
 $= -x[\text{class}] + \log(\sum_j \exp(x[j]))$

Optimizer

Create optimizer

```
optimizer = tf.train.AdamOptimizer(learning_rate)
```

Two methods exists

1. compute gradients and update model

```
grads_vars = optimizer.compute_gradients(loss)  
train_op = optimizer.apply_gradients(grads_vars)
```

2. Update model with one function

```
train_op = optimizer.minimize(loss)
```

* other optimizers?

- `tf.train.GradientDescentOptimizer()`
- `tf.train.RMSPropOptimizer()`
- `tf.train.AdagradOptimizer()`

Now, we can train model with following line

```
sess = tf.Session()  
sess.run(fetches=train_op, feed_dict=inputs)
```

Misc

How to show the variables defined in the network (graph)?

```
for var in tf.global_variables():  
    print(var)
```

* global_variables() returns the variables that are created by tf.Variable() or tf.get_variable()

* We can show operations using tf.get_default_graph().get_operations()

```
Tensor("conv1/weights/read:0", shape=(5, 5, 3, 64), dtype=float32)  
Tensor("conv1/biases/read:0", shape=(64,), dtype=float32)  
Tensor("conv2/weights/read:0", shape=(5, 5, 64, 64), dtype=float32)  
Tensor("conv2/biases/read:0", shape=(64,), dtype=float32)  
Tensor("fc3/weights/read:0", shape=(1600, 384), dtype=float32)  
Tensor("fc3/biases/read:0", shape=(384,), dtype=float32)  
Tensor("fc4/weights/read:0", shape=(384, 192), dtype=float32)  
Tensor("fc4/biases/read:0", shape=(192,), dtype=float32)  
Tensor("fc5/weights/read:0", shape=(192, 10), dtype=float32)  
Tensor("fc5/biases/read:0", shape=(10,), dtype=float32)
```

How to access weight parameters?

```
sess = tf.Session()  
g = tf.get_default_graph()  
w_tensor = g.get_tensor_by_name("fc5/weights/read:0")  
w = w_tensor.eval(session=sess)  
print(w.shape)  
print(w[0, :10])
```

```
# create session  
# get the current graph model  
# obtain a tensor variable using the name  
# get the tensor by running session
```

```
(192, 10)  
[ 0.00417436  0.01035099  0.00198316  0.00643856  0.00888029  0.00903297  
 -0.0122347  -0.01635222 -0.01234471  0.00342033]  
[ 0.00417436  0.01035099  0.00198316  0.00643856  0.00888029  0.00903297  
 -0.0122347  -0.01635222 -0.01234471  0.00342033]
```

Check the code `image_classification.ipynb`

Train the model and evaluate it

Exercises

1. Use residual block instead of simple convolutional layer

a. <https://arxiv.org/pdf/1512.03385.pdf>

b. <https://arxiv.org/pdf/1603.05027.pdf>

